

HW2

Zayd Abdalla, Areeya Aksornpan

3/15/2021

Problem 1: visualization

```
read_data <- function(df) {
  path <- paste("https://raw.githubusercontent.com/jgscott/EC0395M/master/data/",
               df, sep = "")
  df <- read_csv(path)
  return(df)
}

capmetro <- read_data("capmetro_UT.csv") %>%
  mutate(day_of_week = factor(day_of_week,
                              levels = c("Mon", "Tue", "Wed", "Thu",
                                           "Fri", "Sat", "Sun")),
         month = factor(month, levels = c("Sep", "Oct", "Nov")))

##
## -- Column specification -----
## cols(
##   timestamp = col_datetime(format = ""),
##   boarding = col_double(),
##   alighting = col_double(),
##   day_of_week = col_character(),
##   temperature = col_double(),
##   hour_of_day = col_double(),
##   month = col_character(),
##   weekend = col_character()
## )

Figure1 <-
  capmetro %>%
  group_by(hour_of_day, day_of_week, month) %>%
  mutate(avg_boarding = mean(boarding)) %>%
  ungroup() %>%
  ggplot() +
  geom_line(aes(x = hour_of_day, y = avg_boarding, color = month)) +
  scale_x_continuous(expand = c(0,0), limits = c(0, 24),
                    breaks = seq(10, 20, 5)) +
  scale_y_continuous(expand = c(0,0), limits = c(0, 200)) +
  scale_color_ft("Month") +
  facet_wrap(. ~ day_of_week, scales = "free") +
  labs(x = "Hour of day", y = "Average boarding",
```

```

title = "Average bus ridership around UT",
subtitle = "Tracked by Optical Scanner",
caption = "Source: Capital Metro") +
theme_ipsum(grid = "XY", axis = "xy")

```

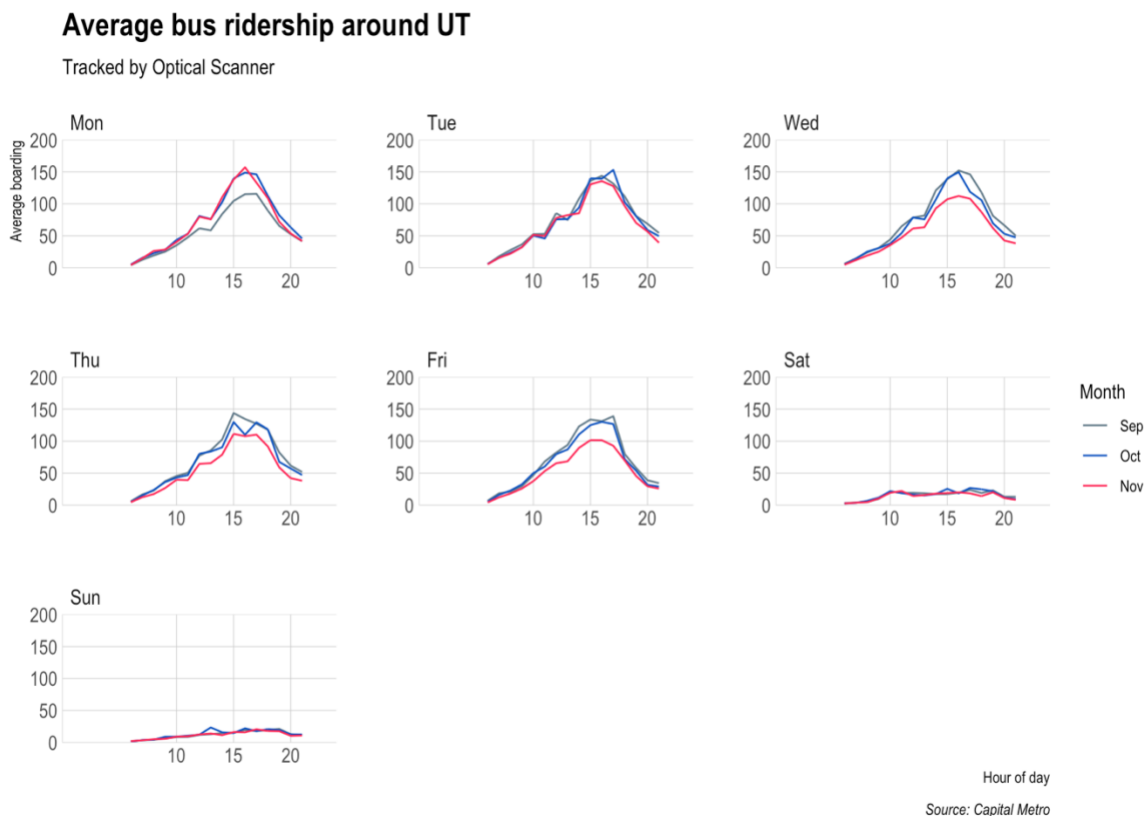


Figure1 illustrates the average CapMetro bus boardings-tracked by Optical Scanner-on weekdays in September, October, and November. The hour of peak boarding appears broadly similar across days, generally peaking around the 17th hour (5pm). This result is intuitive since most people finish school/work around that time. However, weekends tend to not peak in average bus boardings around certain hours as sharply, which supports my intuition that these trends are indicating work commutes. One guess for the decline in average boardings on Mondays in September is that the first Monday of September is Labor Day. Since Labor Day is a holiday, work commutes that day will decline relative to other Mondays, so the average bus boardings in September declines. One guess for the decline in average boardings on Weds/Thurs/Fri in November are because many schools and occupations go on break after the Tuesday before Thanksgiving, which gives people time off from work, so they are less likely to commute on those days.

```

Figure2 <-
  capmetro %>%
  group_by(timestamp, hour_of_day) %>%
  mutate(avg_boarding = mean(boarding)) %>%

```

```

ggplot() +
  geom_point(aes(x = temperature, y = avg_boarding, color = weekend)) +
  scale_x_continuous(expand = c(0,0), limits = c(30, 100),
                    breaks = seq(40, 100, 20)) +
  scale_y_continuous(expand = c(0,0), limits = c(0, 300)) +
  scale_color_ft() +
  facet_wrap(. ~ hour_of_day, scales = "free") +
  labs(x = "Temperature", y = "Boarding",
       title = "Average bus ridership around UT by temperature",
       subtitle = "Faceted by hour of day",
       caption = "Source: Capital Metro") +
  theme_ipsum(grid = "XY", axis = "xy") +
  theme(legend.title = element_blank())

```

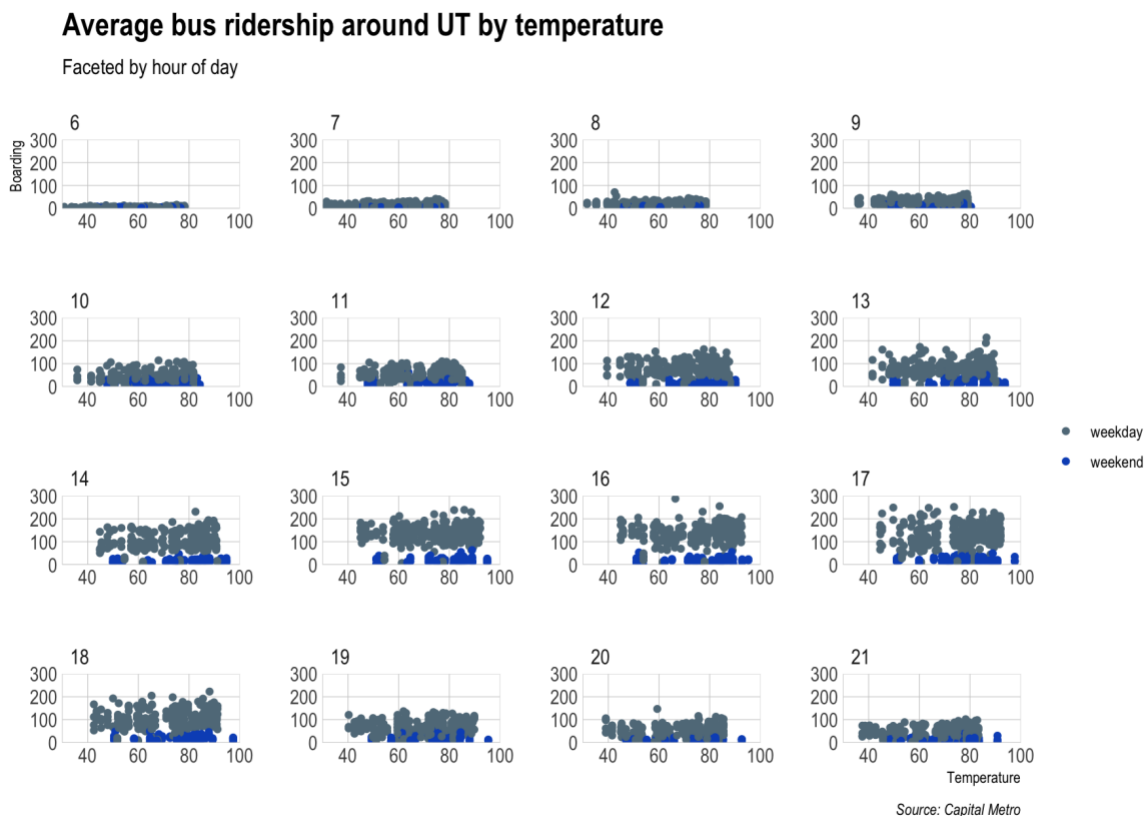


Figure2 shows average ridership, by temperature, which is faceted by hour of the day (6am to 10pm), and averaged by 15-minute increments. Gray indicates weekdays and blue indicates weekends. When we hold hour of day and weekend status constant, temperature does not appear to noticeably change the average ridership of UT students. The changes in bus demand seems to be more related to the time of day since the average boardings at each hour is pretty similar across temperatures.

Problem 2: Saratoga House Prices

```
saratoga <- mosaicData::SaratogaHouses

#create the train/test split.

set.seed(300)

saratoga_split <- initial_split(saratoga, strata = "price", prop = 0.75)
saratoga_train <- training(saratoga_split)
saratoga_test  <- testing(saratoga_split)

dim(saratoga_train)

## [1] 1298  16

dim(saratoga_split)

##   analysis assessment      n      p
##   1298      430    1728    16

#use cross-validation to split training set into k-folds.

# 3 fold cross validation
saratoga_fold <- vfold_cv(saratoga_train, v = 3, repeats = 1, strata = "price")

# Linear and Knn models
lin_mod <-
  linear_reg() %>%
  set_mode("regression") %>%
  set_engine("lm")
lin_mod

## Linear Regression Model Specification (regression)
##
## Computational engine: lm

knn_mod <-
  nearest_neighbor(
    mode = "regression",
    neighbors = tune("K"),
  ) %>%
  set_engine("kknn")
knn_mod

## K-Nearest Neighbor Model Specification (regression)
##
## Main Arguments:
##   neighbors = tune("K")
```

```

##
## Computational engine: kkn

#Use tidymodels to feature engineer: rescaling and standardizing variables
saratoga_wf <-
  workflow() %>%
  add_formula(price ~ .) %>%
  # log price
  step_log(price) %>%
  # mean impute numeric variables
  step_meanimpute(all_numeric(), -all_outcomes()) %>%
  # rescale all numeric variables to lie between 0 and 1
  step_range(all_numeric(), min = 0, max = 1) %>%
  # one-hot
  step_dummy(fuel, centralAir, heating, newConstruction, waterfront, sewer) %
>%
  # remove predictor variables that are almost the same for every entry
  step_nzv(all_predictors())
saratoga_wf

## == Workflow =====
=====
## Preprocessor: Formula
## Model: None
##
## -- Preprocessor -----
-----
## price ~ .

#Fitting LM model

set.seed(400)
lm_rs <-
  saratoga_wf %>%
  add_model(lin_mod) %>%
  fit_resamples(
    resamples = saratoga_fold,
    control = control_resamples(save_pred = TRUE)
  )

#Fitting KNN model

set.seed(400)
# feature engineering
knn_recipe <-
  recipe(price ~ ., data = saratoga_train) %>%
  # log price
  step_log(price) %>%

```

```

# mean impute numeric variables
step_meanimpute(all_numeric(), -all_outcomes()) %>%
# rescale all numeric variables to lie between 0 and 1
step_range(all_numeric(), min = 0, max = 1) %>%
# one-hot
step_dummy(fuel, centralAir, heating, newConstruction, waterfront, sewer) %
>%
# remove predictor variables that are almost the same for every entry
step_nzv(all_predictors())
# workflow
knn_wf <-
  workflow() %>%
  add_model(knn_mod) %>%
  add_recipe(knn_recipe)
# hyperparameter tuning
gridvals <- tibble(K = seq(1, 200))

knn_rs <-
  knn_wf %>%
  tune_grid(
    resamples = saratoga_fold,
    grid = gridvals,
    control = control_resamples(save_pred = TRUE))
knn_rs

## # Tuning results
## # 3-fold cross-validation using stratification
## # A tibble: 3 x 5
##   splits          id   .metrics          .notes          .predictions
##   <list>         <chr> <list>          <list>          <list>
## 1 <split [863/435]> Fold1 <tibble [400 x 5~ <tibble [0 x 1~ <tibble [87,00
0 x 5~
## 2 <split [866/432]> Fold2 <tibble [400 x 5~ <tibble [0 x 1~ <tibble [86,40
0 x 5~
## 3 <split [867/431]> Fold3 <tibble [400 x 5~ <tibble [0 x 1~ <tibble [86,20
0 x 5~

set.seed(400)

# Display only minimum RMSE
knn_min <- knn_rs %>%
  collect_metrics() %>%
  filter(.metric == "rmse") %>%
  filter(mean == min(mean))
knn_min

## # A tibble: 1 x 7
##       K .metric .estimator   mean     n std_err .config
##   <int> <chr>   <chr>       <dbl> <int>   <dbl> <chr>
## 1    28 rmse    standard    0.0667     3 0.00700 Preprocessor1_Model028

```

Evaluate Models

Evaluate Linear Model

```
final_lm_wf <-
  saratoga_wf %>%
  add_model(lin_mod)

lm_fit <-
  final_lm_wf %>%
  last_fit(split = saratoga_split)
lm_fit %>% collect_metrics()

## # A tibble: 2 x 4
##   .metric .estimator .estimate .config
##   <chr>   <chr>      <dbl> <chr>
## 1 rmse    standard    54998. Preprocessor1_Model1
## 2 rsq     standard     0.642 Preprocessor1_Model1

lm_results <-
  lm_fit %>%
  collect_predictions()
# view results
lm_results

## # A tibble: 430 x 5
##   id          .pred .row price .config
##   <chr>      <dbl> <int> <int> <chr>
## 1 train/test split 188781.     8 170000 Preprocessor1_Model1
## 2 train/test split 176225.     9  90000 Preprocessor1_Model1
## 3 train/test split 226408.    11 325000 Preprocessor1_Model1
## 4 train/test split 277084.    26 248800 Preprocessor1_Model1
## 5 train/test split 135004.    27 135000 Preprocessor1_Model1
## 6 train/test split 178201.    30 140000 Preprocessor1_Model1
## 7 train/test split 225475.    32 187000 Preprocessor1_Model1
## 8 train/test split 232186.    36 169900 Preprocessor1_Model1
## 9 train/test split 218866.    37 209900 Preprocessor1_Model1
## 10 train/test split 229481.    38 169900 Preprocessor1_Model1
## # ... with 420 more rows

lm_fit$.workflow[[1]] %>%
  tidy() %>%
  kable(digits = 4, "pipe")
```


term	estimate	std.error	statistic	p.value
(Intercept)	119752.7855	22181.0183	5.3989	0.0000
lotSize	7563.3591	2435.0849	3.1060	0.0019
age	-200.0885	68.0400	-2.9407	0.0033
landValue	0.9023	0.0560	16.1176	0.0000
livingArea	67.8700	5.3229	12.7505	0.0000

pctCollege	-138.4247	182.4636	-0.7586	0.4482
bedrooms	-8262.4653	3008.5782	-2.7463	0.0061
fireplaces	3083.3949	3487.1396	0.8842	0.3767
bathrooms	25595.5289	3998.7661	6.4009	0.0000
rooms	3512.9442	1134.4768	3.0965	0.0020
heatinghot water/steam	-10979.9277	4959.6157	-2.2139	0.0270
heatingelectric	1278.0498	14903.4144	0.0858	0.9317
fuel electric	-12672.5264	14752.0410	-0.8590	0.3905
fuel oil	-41.3979	5894.1796	-0.0070	0.9944
sewerpublic/commercial	1031.3498	4326.6126	0.2384	0.8116
sewer none	-12456.5061	20182.1907	-0.6172	0.5372
waterfrontNo	-136887.7740	17052.8276	-8.0273	0.0000
newConstructionNo	49508.0622	8594.8209	5.7602	0.0000
centralAirNo	-9271.3629	4143.6560	-2.2375	0.0254

LM Graphically

```
lm_results %>%
  ggplot(aes(.pred, price)) +
  geom_abline(lty = 2, color = "black", size = 1) +
  geom_point(alpha = 0.5, color = "dark green") +
  labs(
    title = 'Linear Regression Results',
    x = "Truth",
    y = "Predicted price",
    color = NULL
  ) +
  theme_ipsum()
```


Linear Regression Results



Evaluate KNN Model

```
final_knn_wf <-
  knn_wf %>%
  finalize_workflow(knn_min)
knn_fit <-
  final_knn_wf %>%
  last_fit(split = saratoga_split)
knn_fit %>% collect_metrics()

## # A tibble: 2 x 4
##   .metric .estimator .estimate .config
##   <chr>   <chr>      <dbl> <chr>
## 1 rmse    standard      0.0524 Preprocessor1_Model1
## 2 rsq     standard      0.616  Preprocessor1_Model1
```

predictions

```
knn_results <-
  knn_fit %>%
  collect_predictions()
```

view results


```
knn_results

## # A tibble: 430 x 5
##   id          .pred .row price .config
##   <chr>        <dbl> <int> <dbl> <chr>
## 1 train/test split 0.688     8 0.699 Preprocessor1_Model1
```

```
## 2 train/test split 0.701      9 0.573 Preprocessor1_Model1
## 3 train/test split 0.704     11 0.828 Preprocessor1_Model1
## 4 train/test split 0.756     26 0.775 Preprocessor1_Model1
## 5 train/test split 0.645     27 0.653 Preprocessor1_Model1
## 6 train/test split 0.715     30 0.661 Preprocessor1_Model1
## 7 train/test split 0.731     32 0.718 Preprocessor1_Model1
## 8 train/test split 0.777     36 0.699 Preprocessor1_Model1
## 9 train/test split 0.744     37 0.741 Preprocessor1_Model1
## 10 train/test split 0.755     38 0.699 Preprocessor1_Model1
## # ... with 420 more rows
```

```
# KNN Graphically
```

```
knn_results %>%
  ggplot(aes(.pred, price)) +
  geom_abline(lty = 2, color = "black", size = 1) +
  geom_point(alpha = 0.5, color = "dark green") +
  labs(
    title = 'KNN Regression Results',
    x = "Truth",
    y = "Predicted price",
    color = NULL
  ) +
  theme_ipsum()
```



We built two models-a linear model and a KNN model-to predict the price of houses. The base model appeared to perform quite well, so we decided to tweak it by feature engineering to improve the accuracy. We standardized numeric variables to values

between (0,1), applied a log transformation to the price variable and created dummy variables for all “character” encoded variables. Next, both our linear and KNN regression models were trained on 3 folds without repetition. We gave the KNN model a hyperparameter (neighbors) that was tuned using a tuning grid. Then, these models were fit on out-of-sample data and we found that our linear model clearly outperformed the medium model from class. However, our KNN model heavily outperformed even our improved linear model. This exercise illustrates the capability of KNN models to adapt to non-linearities of the data in order to achieve better fits in predicting pricing of houses.

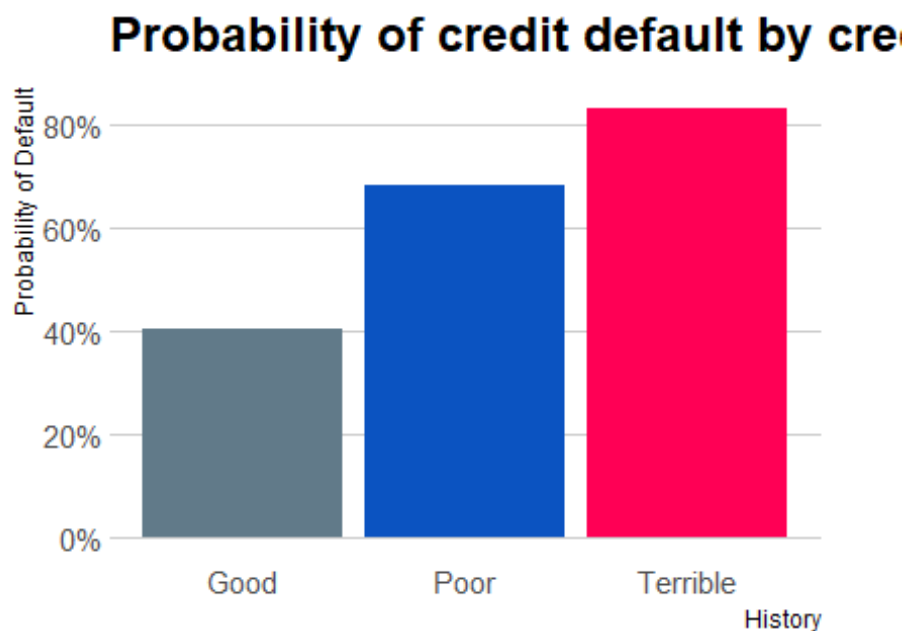
Problem 3: Classification and retrospective sampling

```
german_credit <-  
  read_data("german_credit.csv") %>%  
  select(-1) %>%  
  # Factoring outcomes  
  mutate(Default = as.factor(Default))  
  
## Warning: Missing column names filled in: 'X1' [1]  
  
##  
## -- Column specification -----  
##  
## cols(  
##   .default = col_character(),  
##   X1 = col_double(),  
##   Default = col_double(),  
##   duration = col_double(),  
##   amount = col_double(),  
##   installment = col_double(),  
##   residence = col_double(),  
##   age = col_double(),  
##   cards = col_double(),  
##   liable = col_double(),  
##   rent = col_logical()  
## )  
## i Use `spec()` for the full column specifications.  
  
# Build logistic regression model  
  
german_credit %>%  
  group_by(Default, history) %>%  
  add_tally() %>%  
  rename(num_default = n) %>%  
  distinct(history, num_default) %>%  
  ungroup() %>%  
  group_by(history) %>%  
  mutate(tot_default = sum(num_default),  
         prob_default = (num_default / tot_default) * 100) %>%  
  filter(Default == 0) %>%
```

```

ggplot() +
  geom_col(aes(x = history, y = prob_default,
               fill = history)) +
  scale_y_continuous(labels = function(x) paste0(x, "%")) +
  scale_x_discrete(labels = c("Good", "Poor", "Terrible")) +
  scale_fill_ft() +
  labs(x = "History", y = "Probability of Default",
       title = "Probability of credit default by credit history") +
  theme_ipsum(grid = "Y") +
  theme(legend.title = element_blank(),
        legend.position = "None")

```



```

# Train test
set.seed(395)
german_split <- initial_split(german_credit, strata = "Default", prop = 0.75)
german_train <- training(german_split)
german_test  <- testing(german_split)
# 3 fold cross validation (for speed)
german_fold <- vfold_cv(german_train, v = 3, repeats = 1, strata = "Default")
german_fold

## # 3-fold cross-validation using stratification
## # A tibble: 3 x 2
##   splits          id
##   <list>         <chr>
## 1 <split [500/250]> Fold1

```

```

## 2 <split [500/250]> Fold2
## 3 <split [500/250]> Fold3

# Model engine

log_mod <-
  logistic_reg(penalty = tune(), mixture = 1) %>%
  set_engine("glmnet") %>%
  set_mode("classification")
log_mod

## Logistic Regression Model Specification (classification)
##
## Main Arguments:
##   penalty = tune()
##   mixture = 1
##
## Computational engine: glmnet

# recipe and workflow.

set.seed(350)
# varlist to keep
varlist <- c("Default", "duration", "amount", "installment", "age",
             "history", "purpose", "foreign")
# recipe
log_rec <-
  recipe(Default ~ ., data = german_train) %>%
  # remove vars not in varlist
  step_rm(setdiff(colnames(german_credit), varlist)) %>%
  step_dummy(all_nominal(), -all_outcomes()) %>%
  step_zv(all_predictors()) %>%
  step_normalize(all_predictors())

# workflow
log_wf <-
  workflow() %>%
  add_model(log_mod) %>%
  add_recipe(log_rec)

# Tune grid

log_grid <- tibble(penalty = 10^seq(-4, -1, length.out = 30))

set.seed(350)
log_rs <-
  log_wf %>%
  tune_grid(german_fold,

```


```

      grid = log_grid,
      control = control_grid(save_pred = TRUE),
      metrics = metric_set(roc_auc))
log_rs

## # Tuning results
## # 3-fold cross-validation using stratification
## # A tibble: 3 x 5
##   splits          id    .metrics          .notes          .predictions
##   <list>          <chr> <list>          <list>          <list>
## 1 <split [500/250]> Fold1 <tibble [30 x 5]> <tibble [0 x 1]> <tibble [7,50
0 x 6~
## 2 <split [500/250]> Fold2 <tibble [30 x 5]> <tibble [0 x 1]> <tibble [7,50
0 x 6~
## 3 <split [500/250]> Fold3 <tibble [30 x 5]> <tibble [0 x 1]> <tibble [7,50
0 x 6~

log_rs %>%
  collect_metrics() %>%
  ggplot(aes(x = penalty, y = mean)) +
  geom_point() +
  geom_line() +
  ylab("Area under the ROC Curve") +
  scale_x_log10(labels = scales::label_number()) +
  theme_ipsum()

```



```
top_models <-
  log_rs %>%
  show_best("roc_auc", n = 20) %>%
  arrange(penalty)
top_models %>% kbl(format = "pipe", booktabs = T)
```

penalty	.metric	.estimator	mean	n	std_err	.config
0.0001000	roc_auc	binary	0.7156063	3	0.0286988	Preprocessor1_Model01
0.0001269	roc_auc	binary	0.7156063	3	0.0286988	Preprocessor1_Model02
0.0001610	roc_auc	binary	0.7156063	3	0.0286988	Preprocessor1_Model03
0.0002043	roc_auc	binary	0.7156063	3	0.0286988	Preprocessor1_Model04
0.0002593	roc_auc	binary	0.7156063	3	0.0286988	Preprocessor1_Model05
0.0003290	roc_auc	binary	0.7156063	3	0.0286988	Preprocessor1_Model06
0.0004175	roc_auc	binary	0.7156063	3	0.0286988	Preprocessor1_Model07
0.0005298	roc_auc	binary	0.7156571	3	0.0286530	Preprocessor1_Model08
0.0006723	roc_auc	binary	0.7155810	3	0.0286226	Preprocessor1_Model09
0.0008532	roc_auc	binary	0.7155302	3	0.0287753	Preprocessor1_Model10
0.0010826	roc_auc	binary	0.7155302	3	0.0287700	Preprocessor1_Model11
0.0013738	roc_auc	binary	0.7152762	3	0.0285927	Preprocessor1_Model12
0.0017433	roc_auc	binary	0.7152762	3	0.0286611	Preprocessor1_Model13
0.0022122	roc_auc	binary	0.7147429	3	0.0286857	Preprocessor1_Model14


```

0.0028072 roc_auc binary 0.7142857 3 0.0286611 Preprocessor1_Model15
0.0035622 roc_auc binary 0.7133206 3 0.0283669 Preprocessor1_Model16
0.0045204 roc_auc binary 0.7124571 3 0.0281729 Preprocessor1_Model17
0.0057362 roc_auc binary 0.7108317 3 0.0279406 Preprocessor1_Model18
0.0072790 roc_auc binary 0.7089524 3 0.0271651 Preprocessor1_Model19
0.0092367 roc_auc binary 0.7058794 3 0.0267176 Preprocessor1_Model20
log_rs %>%
  select_best()

## # A tibble: 1 x 2
##   penalty .config
##   <dbl> <chr>
## 1 0.000530 Preprocessor1_Model08

# Model 8 seems to be the best
# Graphically
log_best <-
  log_rs %>%
  collect_metrics() %>%
  arrange(penalty) %>%
  slice(8)
log_auc <-
  log_rs %>%
  collect_predictions(parameters = log_best) %>%
  roc_curve(Default, .pred_0) %>%
  mutate(model = "Logistic Regression")
autoplot(log_auc)

```

```
final_log_wf <-
  log_wf %>%
  finalize_workflow(log_best)
log_fit <-
  final_log_wf %>%
  last_fit(split = german_split)
log_fit %>% collect_metrics()

## # A tibble: 2 x 4
##   .metric .estimator .estimate .config
##   <chr>   <chr>      <dbl> <chr>
## 1 accuracy binary      0.704 Preprocessor1_Model1
## 2 roc_auc  binary      0.721 Preprocessor1_Model1


log_results <-
  log_fit %>%
  collect_predictions()

log_results

## # A tibble: 250 x 7
##   id          .pred_0 .pred_1 .row .pred_class Default .config
##   <chr>         <dbl>  <dbl> <int> <fct>      <fct>   <chr>
## 1 train/test split 0.805  0.195   14 0         1      Preprocessor
## 1_Mod~
## 2 train/test split 0.654  0.346   16 0         1      Preprocessor
## 1_Mod~
```

```
## 3 train/test split 0.842 0.158 22 0 0 Preprocessor
1_Mod~
## 4 train/test split 0.667 0.333 27 0 0 Preprocessor
1_Mod~
## 5 train/test split 0.661 0.339 32 0 0 Preprocessor
1_Mod~
## 6 train/test split 0.917 0.0827 34 0 0 Preprocessor
1_Mod~
## 7 train/test split 0.578 0.422 41 0 0 Preprocessor
1_Mod~
## 8 train/test split 0.749 0.251 42 0 0 Preprocessor
1_Mod~
## 9 train/test split 0.937 0.0627 48 0 0 Preprocessor
1_Mod~
## 10 train/test split 0.803 0.197 49 0 0 Preprocessor
1_Mod~
## # ... with 240 more rows

log_results %>%
  roc_curve(Default, .pred_0) %>%
  autoplot()
```



```
# Confusion matrix
cm <- log_results %>%
  conf_mat(Default, .pred_class)
cm
```

```
##           Truth
## Prediction  0   1
##           0 151  50
##           1  24  25

# Poor sampling
german_credit %>%
  group_by(history) %>%
  tally() %>%
  kbl(format = "pipe")
```

history	n
good	89
poor	618
terrible	293

Our model is accurate roughly 74.4 percent of the time, which is not ideal since our null model that assumes no one will default would be correct 70 percent of the time. We believe the data is not likely ideal for predicting due to the weight of the poorly sampled history variable. Specifically, observe the vast disparity in sampling above.

Problem 4: Children and Hotel Reservations

```
hotels_dev <-
  read_data("hotels_dev.csv") %>%
  mutate(children = as.factor(children))

##
## -- Column specification -----
##
## cols(
##   .default = col_double(),
##   hotel = col_character(),
##   meal = col_character(),
##   market_segment = col_character(),
##   distribution_channel = col_character(),
##   reserved_room_type = col_character(),
##   assigned_room_type = col_character(),
##   deposit_type = col_character(),
##   customer_type = col_character(),
##   required_car_parking_spaces = col_character(),
##   arrival_date = col_date(format = "")
## )
## i Use `spec()` for the full column specifications.

hotels_dev %>%
  count(children) %>%
  mutate(prop = round( n/sum(n), 3)) %>%
  mutate(children = if_else(children == 1, "children", "none")) %>%
  kbl("pipe")
```

children	n	prop
none	41365	0.919
children	3635	0.081

Children only make up about 8% of the sample

```
hotel_splits <- initial_split(hotels_dev, strata = children)
hotel_train <- training(hotel_splits)
hotel_test <- testing(hotel_splits)

train_val_set <- validation_split(hotel_train, strata = children, prop = 0.8)
```

Proportion of children in train/test

train

```
hotel_train %>%
  count(children) %>%
  mutate(prop = round( n/sum(n), 3)) %>%
  mutate(children = if_else(children == 1, "children", "none")) %>%
  kbl("pipe")
```

children	n	prop
none	31033	0.919
children	2717	0.081

test

```
hotel_test %>%
  count(children) %>%
  mutate(prop = round( n/sum(n), 3)) %>%
  mutate(children = if_else(children == 1, "children", "none")) %>%
  kbl("pipe")
```

children	n	prop
none	10332	0.918
children	918	0.082

Both the splits are similar in proportion for children and no children.

Baseline models

penalized logistic regression model

```
log_mod_base1 <-
  logistic_reg(penalty = tune(), mixture = 1) %>%
  set_engine("glmnet")
```

Preprocess recipe

```
log_mod_base1_recipe <-
  recipe(children ~ market_segment + adults + customer_type + is_repeated_guest,
```

```

      data = hotel_train) %>%
    step_dummy(all_nominal(), -all_outcomes()) %>%
    step_zv(all_predictors()) %>%
    step_normalize(all_predictors())

# Preprocess
log_mod_base1_wrkflow <-
  workflow() %>%
  add_model(log_mod_base1) %>%
  add_recipe(log_mod_base1_recipe)

# Tune hyperparameter

lr_reg_grid <- tibble(penalty = 10^seq(-4, -1, length.out = 30))
log_base1_res <-
  log_mod_base1_wrkflow %>%
  tune_grid(train_val_set,
            grid = lr_reg_grid,
            control = control_grid(save_pred = T),
            metrics = metric_set(roc_auc))


# select the best model
log_base1_res %>%
  select_best()

## # A tibble: 1 x 2
##   penalty .config
##   <dbl> <chr>
## 1 0.00924 Preprocessor1_Model20

best_mod_base1 <-
  log_base1_res %>%
  collect_metrics() %>%
  slice(20)

# roc curve
log_base1_res %>%
  collect_predictions(parameters = best_mod_base1) %>%
  roc_curve(children, .pred_0) %>%
  mutate(model = "Logistic Regression") %>%
  autoplot()

```



This is awful
Confusion matrix

```
param_final <-
  log_base1_res %>%
  select_best(metric = "roc_auc")
log_mod_base1_wrkflow <-
  log_mod_base1_wrkflow %>%
  finalize_workflow(param_final)
base1_fit <-
  log_mod_base1_wrkflow %>%
  last_fit(hotel_splits)
base1_pred <-
  base1_fit %>%
  collect_predictions()
base1_pred %>%
  conf_mat(truth = children, estimate = .pred_class)
```

```
##           Truth
## Prediction    0    1
##           0 10332  918
##           1     0    0
```

Baseline 2

Preprocess recipe

```
holidays <- c("AllSouls", "AshWednesday", "ChristmasEve", "Easter",
```

```

      "ChristmasDay", "GoodFriday", "NewYearsDay", "PalmSunday")
log_mod_base2_recipe <-
  recipe(children ~ .,
    data = hotel_train) %>%

  step_date(arrival_date) %>%
  step_holiday(arrival_date, holidays = holidays) %>%
  step_rm(arrival_date) %>%
  step_dummy(all_nominal(), -all_outcomes()) %>%


  step_zv(all_predictors()) %>%
  step_normalize(all_predictors())
# Preprocess
log_mod_base2_wrkflow <-
  workflow() %>%
  add_model(log_mod_base1) %>%
  add_recipe(log_mod_base2_recipe)

# Tune Hyperparameter

log_base2_res <-
  log_mod_base2_wrkflow %>%
  tune_grid(train_val_set,
    grid = lr_reg_grid,
    control = control_grid(save_pred = T),
    metrics = metric_set(roc_auc))

log_base2_res %>%
  collect_metrics() %>%
  ggplot(aes(x = penalty, y = mean)) +
  geom_point() +
  geom_line() +
  ylab("Area under the ROC Curve") +
  scale_x_log10(labels = scales::label_number()) +
  theme_clean()

```



```
# select the best model
```

```
log_base2_res %>%
  select_best()
```

```
## # A tibble: 1 x 2
```

```
##   penalty .config
```

```
##   <dbl> <chr>
```

```
## 1 0.000530 Preprocessor1_Model08
```

```
best_mod_base2 <-
```

```
  log_base1_res %>%
```

```
  collect_metrics() %>%
```

```
  slice(8)
```

```
# roc curve
```


```
log_base2_res %>%
```

```
  collect_predictions(parameters = best_mod_base2) %>%
```

```
  roc_curve(children, .pred_0) %>%
```

```
  mutate(model = "Logistic Regression") %>%
```

```
  autoplot()
```

Not great, but better than baseline1 at least

```
param_final <-
  log_base2_res %>%
  select_best(metric = "roc_auc")
log_mod_base2_wrkflow <-
  log_mod_base2_wrkflow %>%
  finalize_workflow(param_final)
base2_fit <-
  log_mod_base2_wrkflow %>%
  last_fit(hotel_splits)
base2_pred <-
  base2_fit %>%
  collect_predictions()
base2_pred %>%
  conf_mat(truth = children, estimate = .pred_class)
```

```
##           Truth
## Prediction    0    1
##           0 10218  620
##           1   114  298
```

Better

Best Linear Model

```

set.seed(400)
# train/test

hotel_splits2 <- initial_split(hotels_dev, strata = children)
hotel_train2 <- training(hotel_splits2)
hotel_test2 <- testing(hotel_splits2)

# cross-val folds
hotel_cv <- vfold_cv(hotel_train2, v = 10, repeats = 1, strata = children)

# validation set
hotels_val <- read_data("hotels_val.csv") %>%
  mutate(children = as.factor(children))

##
## -- Column specification -----
##
## cols(
##   .default = col_double(),
##   hotel = col_character(),
##   meal = col_character(),
##   market_segment = col_character(),
##   distribution_channel = col_character(),
##   reserved_room_type = col_character(),
##   assigned_room_type = col_character(),
##   deposit_type = col_character(),
##   customer_type = col_character(),
##   required_car_parking_spaces = col_character(),
##   arrival_date = col_date(format = "")
## )
## i Use `spec()` for the full column specifications.

log_mod_rec <-
  recipe(children ~ .,
    data = hotel_train2) %>%

    step_date(arrival_date) %>%
    step_holiday(arrival_date, holidays = timeDate::listHolidays("US")) %>%
    step_rm(arrival_date) %>%
    step_dummy(all_nominal(), -all_outcomes()) %>%
    step_zv(all_predictors()) %>%
    step_normalize(all_predictors())
log_mod <-
  logistic_reg(penalty = tune(), mixture = 1) %>%
  set_engine("glmnet")

# Preprocess
log_mod_wrkflow <-
  workflow() %>%

```

```

add_model(log_mod) %>%
add_recipe(log_mod_rec)

# fit to validation set
hotel_res <-
  log_mod_wrkflow %>%
  tune_grid(grid = lr_reg_grid,
            resamples = hotel_cv,
            control = control_grid(save_pred = T),
            metrics = metric_set(roc_auc))
top_models <-
  hotel_res %>%
  show_best("roc_auc", n = 15) %>%
  arrange(penalty)
hotel_best <-
  hotel_res %>%
  collect_metrics() %>%
  arrange(penalty) %>%
  slice(8)

hotel_best %>% kbl("pipe")


```

			.estimator				
penalty	.metric	r	mean	n	std_err	.config	
0.001373	roc_auc	binary	0.874951	1	0.003565	Preprocessor1_Model0	
8	c		5	0	9	8	

```

# roc curve
hotel_res %>%
  collect_predictions(parameters = hotel_best) %>%
  roc_curve(children, .pred_0) %>%
  mutate(model = "Logistic Regression") %>%
  autoplot()

```



Model Validation 1 using new dataset


```

trained_wf <-
  log_mod_wrkflow %>%
  finalize_workflow(hotel_best) %>%
  fit(hotels_dev)
hotel_preds <-
  trained_wf %>%
  predict(hotels_val) %>%
  bind_cols(hotels_val %>% select(children))
hotel_final_pred <-
  trained_wf %>%
  predict(hotels_val, type = "prob") %>%
  bind_cols(hotel_preds)
hotel_final_pred %>%
  conf_mat(truth = children, .pred_class)

##           Truth
## Prediction    0    1
##           0 4542  261
##           1   55  141

hotel_final_pred %>%
  roc_curve(children, .pred_0) %>%
  mutate(model = "Logistic Regression") %>%
  autoplot()

```



Model Validation 2

create v-folds

```
hotel_folds <- vfold_cv(hotels_val, v = 20)
```

```
hotel_fold_fit <-
  trained_wf %>%
  fit_resamples(
    resamples = hotel_folds,
    control = control_resamples(save_pred = TRUE)
  )
```

predicted probabilities

```
pred_sums <- list()
for (i in 1:20) {
  pred_sums <-
    hotel_fold_fit$.predictions[[i]] %>%
    summarize(sum_pred = sum(as.numeric(.pred_class))) %>%
    pull(sum_pred) %>%
    append(pred_sums)
}
```

actual probabilities

```
actual_sums <- list()
```

```

for (i in 1:20) {
  actual_sums <-
    hotel_fold_fit$.predictions[[i]] %>%
    summarize(sum_actual = sum(as.numeric(children))) %>%
    pull(sum_actual) %>%
    append(actual_sums)
}
# colnames
names <- tibble("Folds" = c("Actual", "Predicted"))

probs <-
  as_tibble(actual_sums, .name_repair = "unique") %>%
  rbind(as_tibble(pred_sums, .name_repair = "unique"))

## New names:
## * `` -> ...1
## * `` -> ...2
## * `` -> ...3
## * `` -> ...4
## * `` -> ...5
## * ...

## New names:
## * `` -> ...1
## * `` -> ...2
## * `` -> ...3
## * `` -> ...4
## * `` -> ...5
## * ...

# Table
cbind(names, probs) %>%
  kable(col.names =
    append("Folds", make.unique(c("Folds", rep("v", 21))), sep = "")[3:2
2]),
    caption = "Sum of probabilities")

```

Sum of probabilities

Folds

v1

v2

v3

v4

v5

v6

v7

v8

v9

v10

v11

v12

v13

v14

v15

v16

v17

v18

v19

v20

Actual

268

273

273

266

265

267

275

268

271

271

273

269

272

263

273

276

269

275

266

268

Predicted

255

260

262

255

256

261

264

258

260

266

263

262

259

256

261

259

260

260


260

262


```

tibble(fold = seq(1, 20, 1),
       Actual = unlist(actual_sums),
       Prediction = unlist(pred_sums)) %>%
  pivot_longer(!fold, names_to = "names", values_to = "vals") %>%
  ggplot() +
  geom_point(aes(x = fold, y = vals, color = names)) +
  labs(x = "Fold", y = "Sum of probabilities") +
  scale_x_continuous(breaks = seq(1, 20, 1)) +
  scale_color_brewer(palette = "Set1") +
  theme_clean() +
  theme(legend.title = element_blank())

```



```

mean_err <- sum(probs[1,] - probs[2,]) / 20
glue::glue("The average mean error is {mean_err}")

```

```
## The average mean error is 10.1
```

The sum of probabilities is fairly similar across folds.